

Final Project

Problem Description:

Bisection and divide-and-conquer: Implement two codes for computing eigenvalues of a symmetric tridiagonal matrix A :

- (a) Method of bisection: Find all the eigenvalues in a given interval $[a, b]$, using a standard bisection method and Sturm sequences.
- (b) Divide-and-conquer: Find all the eigenvalues of A , by solving the secular equation using Newton's method.

Run tests and timings that verify the expected computational complexities of your codes, and compare the accuracy with MATLAB's `eig` function.

Code: Included in the submission, not in this report.

Bisection:

a: High level idea:

Given the nice interlace property of eigenvalues of symmetric diagonal matrix, we can compute the Sturm sequence with any shift S of a matrix in linear time, and use the number of sign changes in the sequence to determine the number of eigenvalues between $(-\infty, S)$.

With this technique in hand, we can quickly find the number of eigenvalues between $[a, b]$, which is already of great value. Furthermore, since we can generally find number of eigenvalues between any interval, we can bisect the given interval $[a, b]$ and find the exact eigenvalues by just bisecting subintervals until subintervals (that contains exactly one eigenvalue) are so small that we can treat any point in the interval as an eigenvalue of A .

- Input: A , a , b , tolerance
- Return: eigenvalues of A in interval $[a, b]$, with relative error under tolerance
- precondition: A is symmetric, tridiagonal. The eigenvalues of A must be unique in context of tolerance. (in other words, the method breaks down with huge tolerance, since the algorithm will stop when there are still several eigenvalues in the interval.)

1. Check if b is an eigenvalue of A .

2. Bisect $[a, b]$ into $[a, (a+b)/2]$ and $[(a+b)/2, b]$, the eigenvalues in $[a, b]$ are just the union of eigenvalues of two subintervals.

3. Find eigenvalues of subintervals recursively by repeating 2, until subinterval is small enough.

b: Choices I have made:

- When calculating the Sturms sequence with large matrices and large shift, the numbers will actually overflow or underflow. So instead of calculating the exact sequence, I scale some terms if they appears too large. I did not implement the underflow case, as I think we can pre-scale the matrix to convert that into an overflow problem, which is much easier to handle.
- The reason why underflow is hard for me is that, in my calculations of Sturms, if the last term of the sequence is really close to zero, I have a good reason to believe the shift is almost an eigenvalue. Since there seems to be a cumulated precision error, when matrix size is large, the algorithm is prone to error when trying to find sign changes near zero. Thus, I will recalculate the Sturm

sequence with a slightly larger shift to avoid this problem.

- The way to calculate the number of sign changes follows closely to the convention introduced in Trefethen/Bau.
- To Avoid dynamic memory allocation, I pre-allocated memory for eigenvalues. This can be buggy if the precondition do not met and my algorithm cannot find enough number of eigenvalues in the subinterval. However, if precondition is satisfied, no error will occur.
- tolerance is user-chosen. The algorithm will break if tolerance is smaller than machine error.
- Sub intervals that has 0 eigenvalues in it will not be examined.
- When bisecting $[a,b]$, I will calculate the number of eigenvalues from $(-\infty, a), (-\infty, \frac{a+b}{2}), (-\infty, b)$. Note that, these values can be reused when bisecting $[a, \frac{a+b}{2}), [\frac{a+b}{2}, b)$. Thus, I will pass these values into smaller function calls, which nearly doubles the speed.

d: Test Problem, Results and Analysis:

1. Random Normal A, find all eigenvalues.

```

1  m = 512; %4,8,16,32,64,128,256,512
2  Error=0;
3  for i=1:10
4      A = randn(m);
5      B = hess(A'*A)*10^40; % huge numbers
6      %B = hess(A'*A);      % normal numbers
7
8      E = eig(B);
9      %ANS = bisection(B,0,10000,10e-14,-1,-1);% normal numbers
10     ANS = bisection(B,0,10000*10^40,10e-14,-1,-1);% huge numbers
11     if size(ANS,1)~=m
12         fprintf('BAD!\n');
13     end
14     Error= Error + norm(E(1:size(ANS,1))-ANS);
15 end
16 Error=Error/10

```

Results: See next page. I ran each problem ten times to average out some outliers and included the total errors. I also included the number of function calls and running time of bisection and Sturm.

Analysis:

The number of function calls for each eigenvalue do not vary with scaling of both numbers and size of the problem, which is expected. In fact, with problem size being larger, the average number of functions calls for each eigenvalue goes down. However, huge numbers requires slightly more time because I have to scale more frequently to keep the Sturm sequence small.

The relative error is the same for different scalings. (Multiplying by 10^{40} gives 10^{40} bigger errors) The time here is quadratic with the problem (matrix) size. Based on the number of function calls, as size of the problem double, we call double amount of Sturm, which is a $O(m)$ function. This means the implementation is $O(m^2)$ indeed.

TABLE 1. Test1

bisection 10 times.							
Scale	Size(m)	Error	# bisection called/time		# Sturm called/time(s)		# of Sturm per eigenvalue
Normal	512	3.60E-10	199911	4.218	194815	40.465	38.04980469
Normal	256	1.21E-10	102533	1.836	99996	10.492	39.0609375
Normal	128	4.48E-11	52579	0.819	51329	2.756	40.10078125
Normal	64	1.63E-11	26952	0.389	26332	0.776	41.14375
Normal	32	5.30E-12	13837	0.211	13545	0.285	42.328125
Normal	16	1.98E-12	7120	0.128	6983	0.117	43.64375
Normal	8	6.37E-13	3705	0.139	3645	0.065	45.5625
Normal	4	2.13E-13	1935	0.06	1929	0.028	48.225
Huge	512	3.55E+30	199852	4.171	194753	45.918	38.03769531
Huge	256	1.23E+30	102575	1.546	100061	11.339	39.08632813
Huge	128	4.52E+29	52590	0.662	51330	2.974	40.1015625
Huge	64	1.51E+29	26980	0.382	26361	0.909	41.1890625
Huge	32	5.01E+28	13851	0.291	13551	0.356	42.346875
Huge	16	1.92E+28	7118	0.152	6979	0.115	43.61875
Huge	8	4.84E+27	3700	0.106	3640	0.072	45.5
Huge	4	1.95E+27	1935	0.039	1915	0.04	47.875
* Everything multiplied by 10^{40}							

2. Random Normal A, find eigenvalues in a range.

```

1 m = 64; %4,8,16,32,64,128,256,512
2 Error=0;
3 a = 1;
4 b = 512; %2, 4,8,16,32,64,128,256
5 Eigenvalues_found =0;
6 for i=1:10
7     A = randn(m);
8     B = hess(A'*A); % normal numbers
9     E = eig(B);
10    ANS = bisection(B,a,b,10e-14,-1,-1); % normal numbers
11
12    E_in_range=zeros(0,1);
13    for j=1:m
14        if (E(j)>=a && E(j)<=b)
15            E_in_range(end+1,1)=E(j);
16        end
17    end
18
19    if size(ANS,1)~=size(E_in_range,1)
20        fprintf('BAD! %d %d \n',size(ANS,1),size(E_in_range,1));
21        ANS
22        E_in_range
23    end
24    Eigenvalues_found = Eigenvalues_found + size(ANS,1);
25    Error= Error + norm(E_in_range-ANS);
26 end
27 Error=Error/10
28 Eigenvalues_found

```

TABLE 2. Test2, matrix is 64 by 64

bisection 10 times						
Range	Eig found	Total error	Error per eigenvalue	# Sturm called/time(s)		# of Sturm per eigenvalue
[1,2]	21	4.05E-14	1.93E-15	911	0.06	43.38095238
[1,4]	51	1.12E-13	2.19E-15	2174	0.117	42.62745098
[1,8]	93	2.39E-13	2.57E-15	3928	0.175	42.23655914
[1,16]	154	6.46E-13	4.19E-15	6453	0.251	41.9025974
[1,32]	230	1.58E-12	6.85E-15	9558	0.276	41.55652174
[1,64]	339	3.63E-12	1.07E-14	13958	0.446	41.1740413
[1,128]	468	8.12E-12	1.73E-14	19146	0.633	40.91025641
[1,256]	584	1.40E-11	2.39E-14	23808	0.705	40.76712329

Analysis:

As we can see, the error per eigenvalue is invariant with the given parameter a and b (in the sense of relative error, for larger eigenvalues, we do expect slightly larger errors). Furthermore, the number of Sturm called for each eigenvalue is around 40. Given that our input relative tolerance is $10^{-13} \approx 2^{-43}$, the number of function calls are slightly better than expected.

Divide and Conquer:

a: High level idea of implementation:

- Using the idea of divide and conquer to solve all eigenvalues of A.
- Return: Sorted eigenvalues of A and corresponding eigenvector matrix Q
- precondition: A is m by m ($m \geq 2$), symmetric, tridiagonal.
- 0. If A is small, solve using ordinary methods.
- 1. If A is not irreducible (There are some zeros in the sub diagonal of A), reduce A.
- 2. Divide: find eigenvalues/eigenvector of upper-half/ lower-half
- 3. MergeSort two eigenvalues returned in 2
- 4. Conquer: Find eigenvalue of $D \pm w^*w$ by solving secular equations.
- 5. Conquer: Given eigenvalue, find eigenvector in $O(m^2)$

b: Choices I have made:

- For the purpose of this project, I will keep dividing A until its size is smaller or equal to 2.
- After I divide A and conquer both top and bottom part, I will use merge-sort to sort their eigenvalues, which is $O(m)$.
- I wrote two files to handle positive and negative case of β .
- I chose 10^{-10} as my tolerance.

c: How to the secular equation formed in the Conquer Phase:

- I basically used a bisection structure, with possible speedup with Newton's method when applicable. I will show the logical sequence of the positive β case, negative β case follows trivially.

Given the eigenvalue D, \vec{w} , sorted D called Dsort.

For each i in 1 to m-1, λ_i will be bounded by D_i and D_{i+1} .

Thus, set **left**= D_i and **right**= D_{i+1} , $\lambda_i = (\text{left and right})/2$.

While True:

—— If **left** and **right** are close enough, return λ_i .

—— Else, find:

—— $f(\lambda_i) = \vec{w}^T(\text{diag}(D) - \lambda_i I)^{-1} \vec{w}$

—— $f'(\lambda_i) = \vec{w}^T(\text{diag}(D) - \lambda_i I)^{-2} \vec{w}$

—— **intercept** = $\lambda_i - f(\lambda_i)/f'(\lambda_i)$

—— If **intercept** is really close to λ_i , we have found a good λ_i , return λ_i

—— If **intercept** is within **left** and **right**, $\lambda_i = \text{intercept}$ and continue

—— Else, do bisection, update **left** and **right** according to $f(\lambda_i)$.

For λ_m , the last eigenvalue, it is only lower bounded by D_m , I used a variation of bisection:

set **left**= D_i and **right**= ∞ , $\lambda_i = 2 * \|\text{left}\|$.

While True:

—— If **left** and **right** are close enough, return λ_i .

—— Else, find:

—— $f(\lambda_i) = \vec{w}^T(\text{diag}(D) - \lambda_i I)^{-1} \vec{w}$

—— If $f(\lambda_i)$ is positive, shift **right** to λ_i and now we can do normal bisection.

—— Else, shift **left** to λ_i and $\lambda_i = 2 * \|\text{left}\|$.

d: Test Problem:

```

1  m = 512;
2  Error=0;
3  for i=1:10
4      A = randn(m);
5      B = hess(A'*A)*10^20;
6      [V1,D1] = eig(B);
7      [V2,D2] = DivideAndConquer(B);
8      Error = (norm(D1-diag(D2)));
9  end
10 Error=Error/10

```

Note that, the unbounded eigenvalue can take huge amount of steps theoretically, but in practice the cost is negligible. Here I use rand normal matrix to prevent this unconsidered factor.

e: Results:

TABLE 3. Normal Scale

Divide and Conquer 10 times (Stable Newton)						
Scale	Size(m)	Error	number of DivideAndConquer called/time		Secular Equations Solved /time(s)	
Normal	512	7.29E-01	5110	12.528	2550	233.829
Normal	256	4.30E-01	2550	1.466	1270	26.213
Normal	128	9.30E-02	1270	0.302	630	3.476
Normal	64	2.02E-10	630	0.152	310	1.257
Normal	32	1.20E-10	310	0.083	150	0.379
Normal	16	1.28E-10	150	0.052	70	0.246

TABLE 4. Huge Scale

Divide and Conquer 10 times (Stable Newton)						
Scale	Size(m)	Error	number of DivideAndConquer called/time		Secular Equations Solved /time(s)	
huge	512	1.28E+20	5110	11.659	2550	188.957
huge	256	1.15E+20	2550	1.407	1270	19.663
huge	128	4.77E+19	1270	0.272	630	2.982
huge	64	2.13E+14	630	0.142	310	0.751
huge	32	7.04E+12	310	0.075	150	0.207
huge	16	3.32E+11	150	0.033	70	0.078
* Everything multiplied by 10^20						

f: Analysis:

Based on the number of function calls, my recursive step looks correct.

Note that, the error do not scale linearly with the matrix size. For the normal case from 64 to 128, there is a huge jump of errors. I have double checked my code and concluded that the errors are resulting from the slightly incorrect eigenvalues and eigenvectors in the divide step. With large matrix size, the probability of this sudden jump occurring is not negligible and account wholly for the increased error. The possible way to fix is to increase tolerance when solving small matrices. The errors for huge numbers have the same pattern and relative error do not scale.

Even with multiple tests, because of the fluctuation of my computer's performance, the relative time between huge and normal cannot be compared. I expect them to be relatively the same. For each set of tests, normal scale, for example, we expect to see a cubic pattern because my implementation not only finds the eigenvalues, but eigenvectors as well. My implementation shows roughly a cubic pattern. Note that even if eigenvalues may only take $O(m^2)$, I have tried for quite a while but did not figure out a way to perform the algorithm without knowing the eigenvectors of sub matrices. I can speedup my implementation by not computing the last eigenvector but that will be a constant time multiplier.

Also, I have tested Stable Newton and bisection only methods. The difference of time is basically a factor of 2, which means, the algorithm is still performing bisection in most cases. (for full Newton I expect a factor of 5 (40 verses $\log(40)$) or so, which did not show up.)